

Weights & Biases Demo

Recyclable Material Classifier

January 8th, 2020

Casey Duncan





About Me

Education



Colorado School of Mines

MS - Mechanical Engineering, Robotics & Automation

Minor - Computer Science

- Main Focus:
 - Robot path planning
 - Robot perception
 - Robot human interaction
- Favorite Topics
 - Computer Vision
 - Reinforcement Learning for path planning

Professional Background

Raytheon Vision Systems

Manufacturing Engineer (2013 - Current)



- Main Responsibilities:
 - Write documentation
 - Train & Assist lab personnel
 - Product-to-Customer delivery process



University of California, Santa Barbara

BS - Mechanical Engineering (2010 - 2014)



Overview

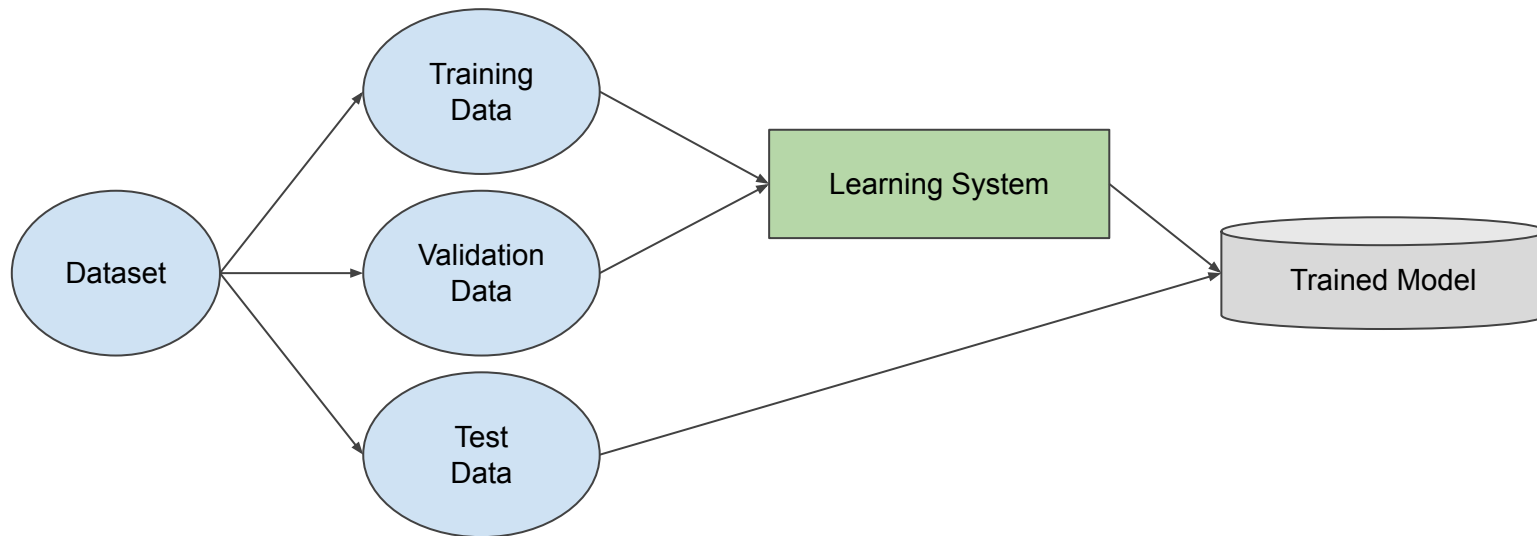
1. Machine Learning / Computer Vision
2. Project Description
3. ML Model Management Approach
 - Previous Approach
 - Weights & Biases
4. Benefits of Weights & Biases
5. Questions



Machine Learning

What is Machine Learning?

- Application of artificial intelligence (AI)
- Allows computer to learn & improve from experience.



Computer Vision

What is Computer Vision?

- Allows computers to gain a high-level understanding from digital photos & videos
- Modern Computer Vision utilizes Machine Learning to increase intelligence

How is it used?

- Facial Recognition
- Autonomous Driving
- Agriculture Crop Selection






Image Classifier

What is an Image Classifier?


- Built using Machine Learning & Computer Vision
- **Input:** Image of an item
- **Output:** Prediction of what item is



Recyclable Material Classifier

What is a Recyclable Material Classifier?

- **Input:** Image of Material
- **Output:** Material Type



Robotic Recyclable Sorter





Motivation

Why Recycling?


- Recycling is hard, even for people
 - Dynamic market, desirable materials constantly shifting
 - Best practices are confusing and locality dependent
- Contribute to the Circular Economy
 - Reduces waste in landfills
 - Reduces overall carbon footprint
- Critical Industry with exceptional returns¹
 - Largest US operator: Waste Management, \$42.2bn Market Cap
 - Up ~300% from 2012

¹ Macrotrends, Waste Management Net Worth 2006 -2019. <https://www.macrotrends.net/stocks/charts/WWM/waste-management/net-worth>



Goal

- Train model with 90% prediction accuracy
 - Leverage transfer learning of proven successful classification models
 - Compare different model performance





Dataset

Collect a dataset of images²

- 6 Classes
 - Paper (594 images)
 - Plastic (482 images)
 - Glass (501 images)
 - Metal (410 images)
 - Cardboard (403 images)
 - Trash (137 images)
- Total Number of Images: 2527 images

Cardboard



Glass



Metal



Paper



Plastic



Trash



² Kaggle, Garbage Classification. <https://www.kaggle.com/asdasdasdasdas/garbage-classification>



Model Optimization

What are Hyperparameters?

- Values used to control the learning process
- Defined prior to training

My Hyperparameters include:

- Data augmentation
- Train / Validation / Test Data Split Ratios
- Training Epoch Quantity
- Learning Rate
- Model Architecture



Hyperparameters

Data Augmentations

- Great for when not enough data to produce robust results
- Augmentation Types:
 - Random Horizontal Flip
 - Random Vertical Flip
 - Random Rotation (Up to $\pm 30^\circ$)
 - Random Artificial Noise
 - Random Gaussian Blur
- Any combination of augmentations can be chosen

Original




Horizontal Flip



Vertical Flip



Rotated 30 Degrees



Noise



Blur

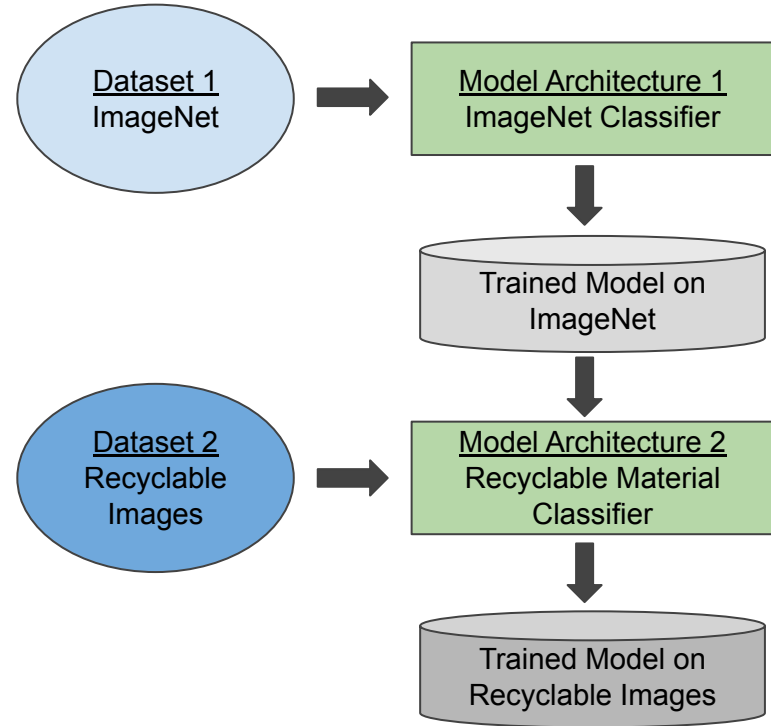




Hyperparameters

Model Architectures

- Pretrained models
 - ResNet 18
 - ResNet 34
 - ResNet 50
- Customized Model
 - ResNet50-ModNet





Model Performance Reporting

Model Performance Metrics

- Training
 - Loss - penalty for a wrong prediction
- Validation
 - Material Accuracy per Epoch
 - Overall Accuracy per Epoch
- Test
 - Material Accuracy
 - Overall Accuracy
 - Confusion Matrix
 - Table displaying predictions vs. actuals
 - **Rows:** Predicted Material
 - **Columns:** Actual Material



Model Performance Reporting

Previous Approach

	A	B	C	D	E	F
1	Model file name:	resnet18_v1				
2						
3	Selected Transforms:					
4	1) Random Horizontal Flip					
5						
6	Training Ratio:	0.68				
7	Validation Ratio:	0.16				
8	Testing Ratio:	0.16				
9						
10	Epoch Quantity:	15				
11						
12	Confusion Matrix					
13	Cardboard	Glass	Metal	Paper	Plastic	Trash
14	57	0	3	1	2	1
15	0	18	29	0	33	2
16	1	1	62	0	2	1
17	3	0	10	54	4	14
18	0	0	11	0	73	1
19	0	0	2	0	0	19
20						
21	Class Accuracy					
22	Cardboard	Glass	Metal	Paper	Plastic	Trash
23	0.890625	0.219512195121951	0.925373134328358	0.635294117647059	0.858823529411765	0.904761904761905
24						
25	Overall Accuracy:	0.70049504950495				



Model Performance Reporting

Weights & Biases - Logging

Logging performance using Weights & Biases is easy!

Example: Logging Test Accuracy

1. Initialize new job before Testing Model

- `run = wandb.init(project="Computer-Vision-Recyclable-Classifiers",
name="resnet18_v1_test",
config=hyperparameters)`

2. Test Model

3. Calculate Test Accuracy

- Save to variable: `test_accuracy`

4. Log Test Accuracy in Weights & Biases Project

- `wandb.log({"Test Accuracy": test_accuracy})`

5. Finish Run


- `run.finish()`



Model Performance Reporting

Weights & Biases - Logging

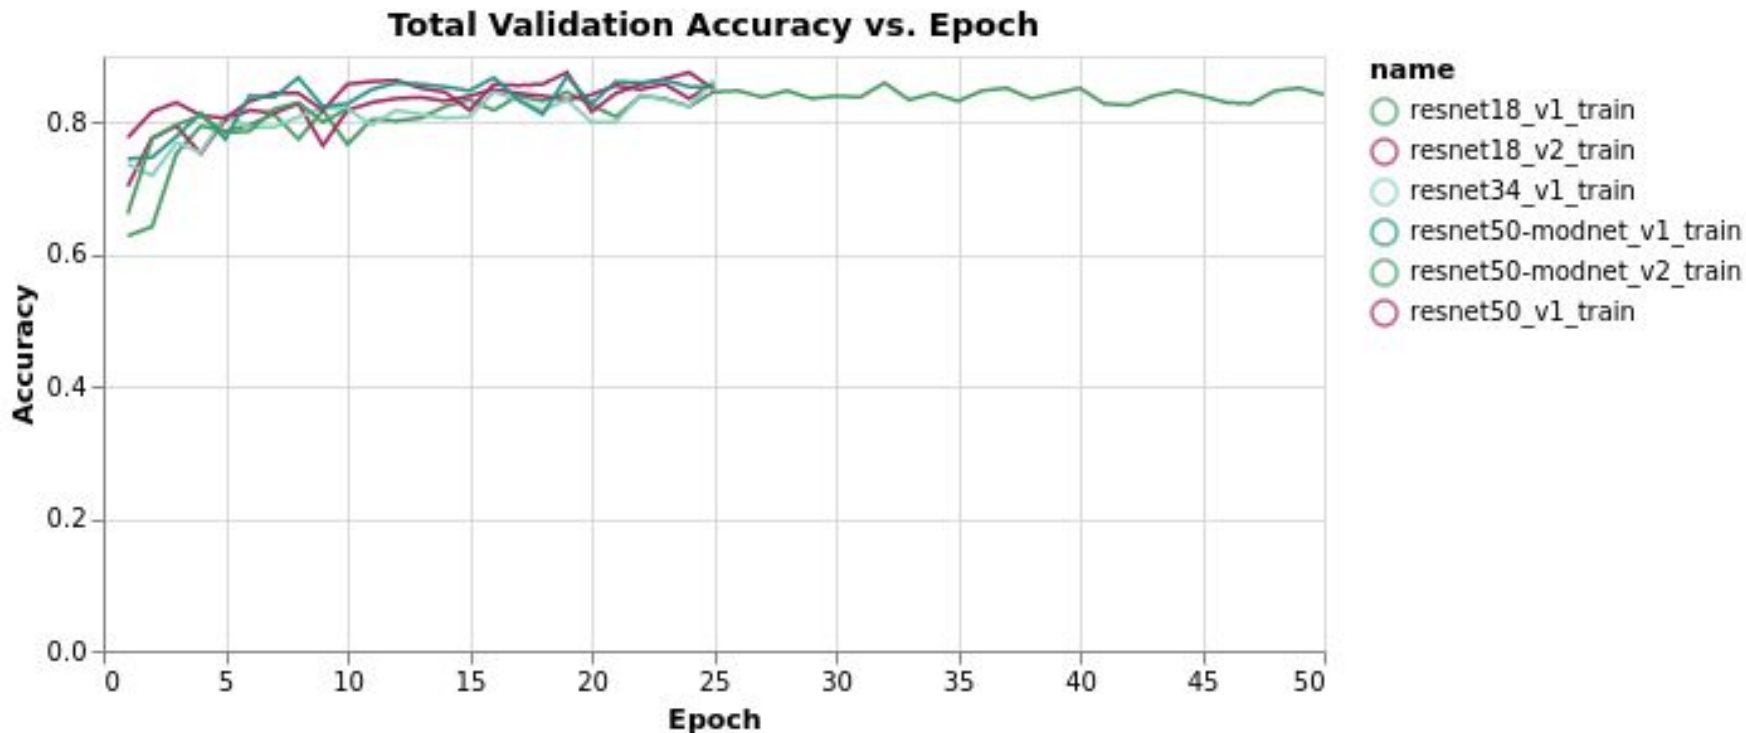
Test Accuracy





Model Performance Reporting

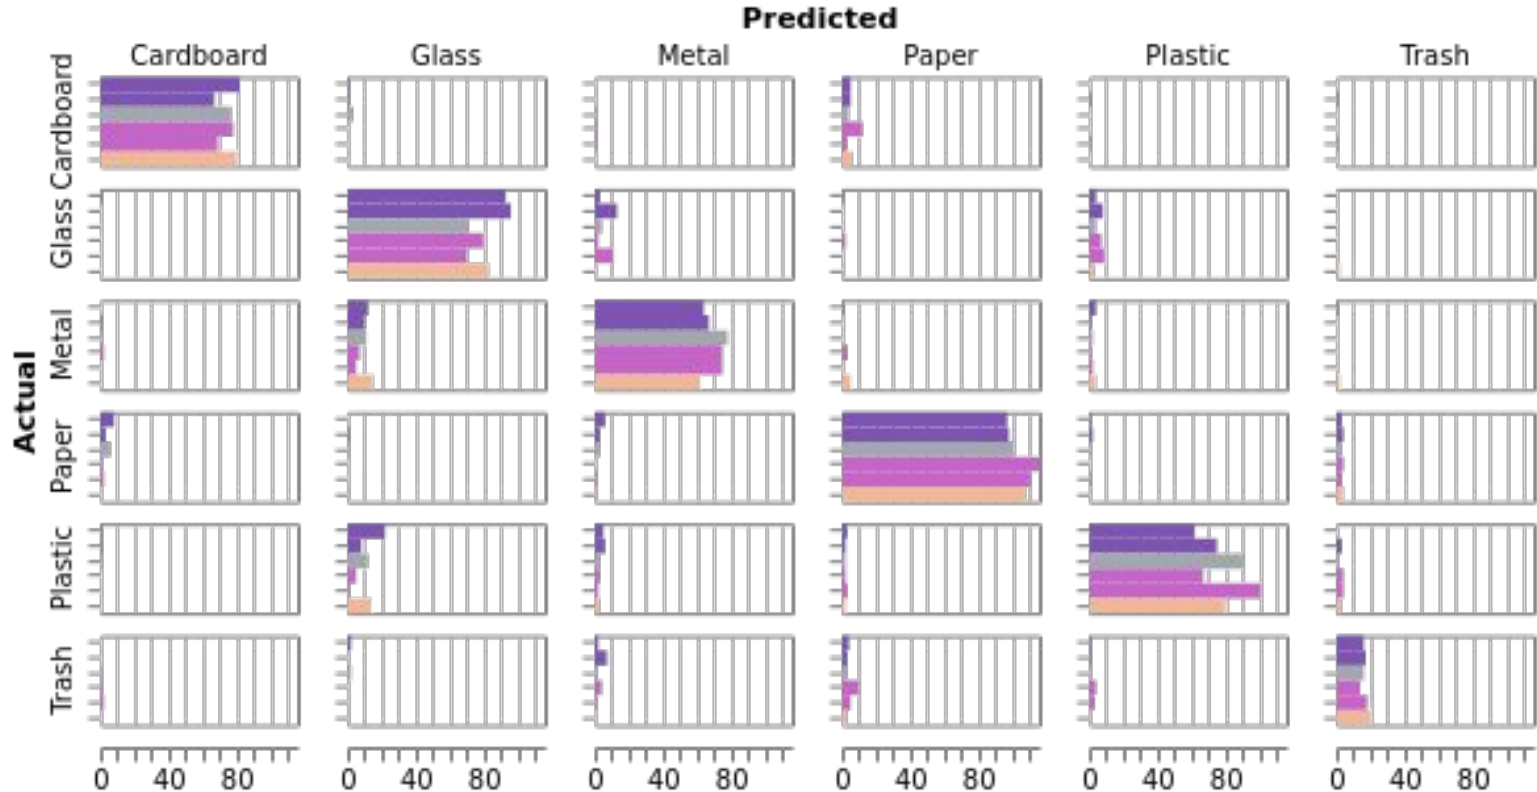
Weights & Biases - Logging





Model Performance Reporting

Weights & Biases - Logging





Saving Files

When training & testing new models, it is important to save:

- Model Files
- Data
 - Raw image data files
 - Training / Validation / Test data files




Saving Files

Previous Approach

Popular method for this is saving to local computer:

- Model Files
 - `PATH = "/Computer-Vision-Recyclable-Classfier/saved_models/resnet18_v1.pth"`
`torch.save(model.state_dict(), PATH)`



- Saving Training / Validation / Test data
 - Did not save



Saving Files

Weights & Biases - Artifacts

Weights & Biases 'Artifacts' feature:

- Saves important files when training & testing new models.
For example:
 - Models
 - Raw Data
 - Data after splitting into Training / Validation / Test
- Creates diagram for how files all tie together



Saving Files

Weights & Biases - Artifacts

Saving File

Example: Saving Model

1. Initialize new job before Training Model

- `run = wandb.init(project="Computer-Vision-Recyclable-Classifier",
name="resnet18_v1_train",
config=hyperparameters)`

2. Train Model

3. Saved to Trained Model file on local computer

4. Log Model Artifact in Weights & Biases Project

- `trained_model_artifact = wandb.Artifact("resnet18_v1_trained",
type="model",
metadata=dict(wandb.config))

trained_model_artifact.add_file(PATH)
run.log_artifact(trained_model_artifact)`

5. Finish Run

- `run.finish()`



Saving Files

Weights & Biases - Artifacts

Type: model

- ▶ resnet50-modnet_v1_trained
- ▶ resnet50-modnet_v1_untrained
- ▼ resnet18_v1_trained
 - v0 latest
- ▶ resnet18_v1_untrained

Overview API Metadata **Files** Graph view

> root

 resnet18_v1_train.pth	44.8MB	
---	--------	---



Saving Files

Weights & Biases - Artifacts

Type: model

▼ resnet18_v1_trained

v0 latest

▶ resnet18_v1_untrained

Overview

API

Metadata

Files

Graph view

Metadata

Q Search

Page 1 of 1

< > ☰

Name	Value
architecture	resnet18
batch_size	8
epoch_qty	10
learn_rate	0.001
test_ratio	0.2
train_ratio	0.6
transform_blur	false
transform_horz	false
transform_noise	false
transform_rot	false



Saving Files

Weights & Biases - Artifacts

Type: split_data

▼ Test_Data_505

v0 latest

▶ Val_Data_505

▶ Train_Data_1517

OverviewAPIMetadataFilesGraph view

> root

Search

<div><div></div><div>Cardboard /</div></div>	1.6MB	88 files
<div><div></div><div>Glass /</div></div>	1.3MB	102 files
<div><div></div><div>Metal /</div></div>	1.4MB	82 files
<div><div></div><div>Paper /</div></div>	2.6MB	115 files
<div><div></div><div>Plastic /</div></div>	1.3MB	93 files
<div><div></div><div>Trash /</div></div>	312.3KB	25 files

Saving Files

Weights & Biases - Artifacts

Type: split_data

▼ Test_Data_505

v0 latest

▶ Val_Data_505

▶ Train_Data_1517

Overview

API

Metadata

Files

Graph view

> root / Glass / glass107.jpg





Saving Files

Weights & Biases - Artifacts

Create Diagram through 'using Artifacts'

Example: Linking Train / Val Data to Model

1. Initialize new job before Training Model

- `run = wandb.init(project="Computer-Vision-Recyclable-Classfier",
name="resnet18_v1_train",
config=hyperparameters)`

2. Use Training & Validation data Artifacts

- `run.use_artifact("Train_Data_180:latest")
run.use_artifact("Val_Data_60:latest")`

3. Train Model

4. Saved to Trained Model file on local computer

5. Log Model Artifact in Weights & Biases Project

6. Finish Run

- `run.finish()`

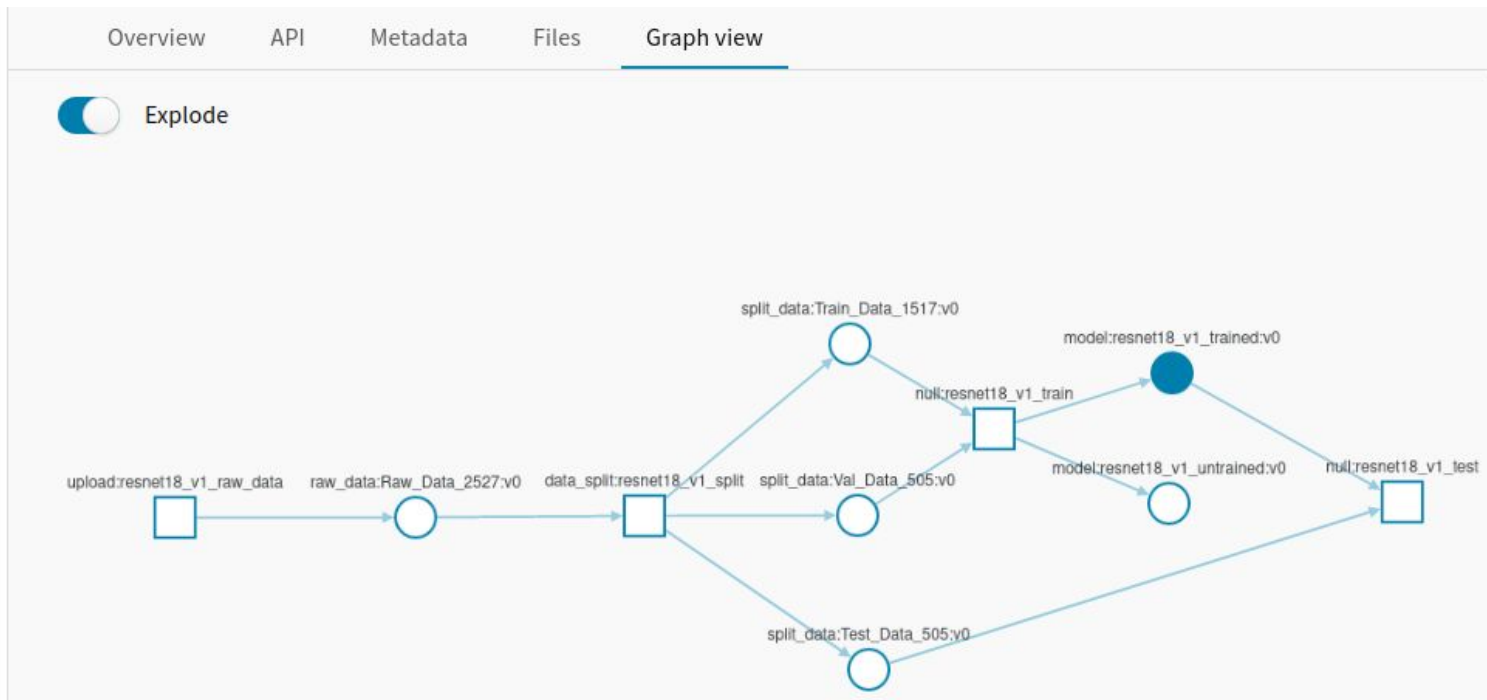


Saving Files

Weights & Biases - Artifacts

Type: model

- ▼ resnet18_v1_trained
 - v0 latest
- ▶ resnet18_v1_untrained





Model Optimization


- How is model performance optimized?
 - Choosing best hyperparameter values
 - There is an art to this due to:
 - Model Architecture
 - Raw Data
 - Project Type



Model Optimization

Previous Approach

- Used excel files to compare different training results





Model Optimization

Weights & Biases - Sweeps

Weights & Biases assists with Model Optimization through the 'Sweeps' feature.

1. Define dictionary of Hyperparameters Defaults at beginning of main file.

- `hyperparameters = dict(train_ratio = 0.6,
val_ratio = 0.2,
test_ratio = 0.2,
epoch_qty = 5,
learn_rate = 0.001,
transform_horz = False,
transform_vert = False,
transform_rot30 = False,
transform_blur = False,
transform_noise = False,
architecture = "resnet18")`

- Tailor code to:

- Use Hyperparameter dictionary
- Accept arguments from command line using hyperparameter variable names.
 - For example, the following should run:

```
$ python main_wandb.py --epoch_qty=8 --transform_horz=True
```



Model Optimization

Weights & Biases - Sweeps

2. Create a YAML file ([sweep.yaml](#)) in base directory containing sweep parameters
 - Add metric Goal, Name, & Target
 - Change hyperparameter ranges
3. Initialize sweep
 - Run the following in the command line

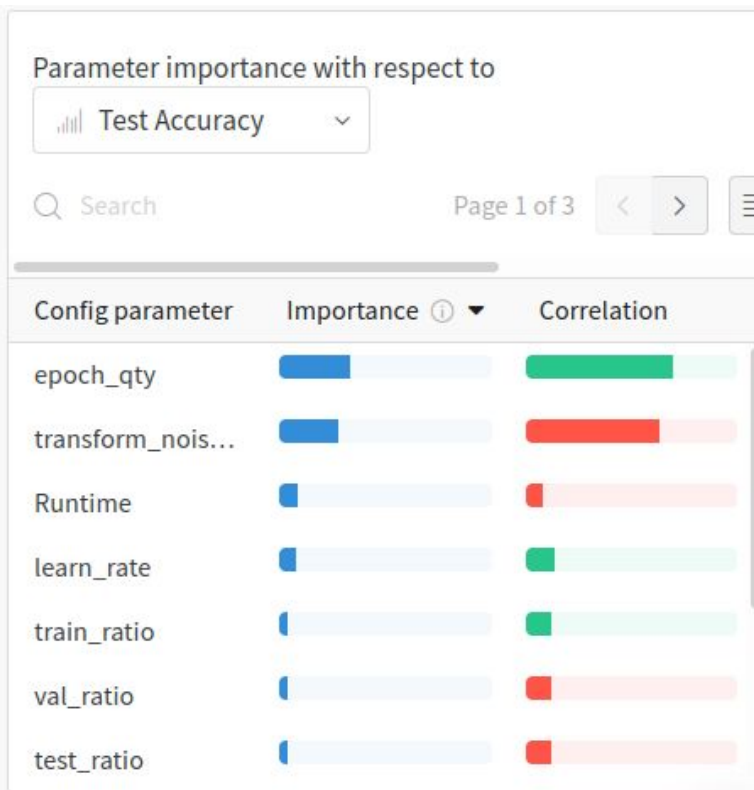

```
$ wandb sweep sweep.yaml
```
4. Run sweep
 - Run sweep command in the command line. Sweep command obtained from step 4.
For example:

```
$ wandb agent caseyduncan/recycling-classifier-demo/7x90t06x
```



Model Optimization

Weights & Biases - Sweeps





Using Plots

Weights & Biases - Reports

Weights & Biases 'Reports' feature lets users easily share model results.

- Simple to build report
- Easily shareable
- Similar to project dashboard

[Example Report](#)



Benefits to Weights & Biases

- Logging Metrics
 - Easily save all performance metrics to W&B project
 - No plotting required
 - Easy to compare different models
- Saving Files
 - Model files
 - Raw & Split Data
 - Creates diagram for how 'Artifacts' connect
- Optimizing Model
 - Intelligently tunes hyperparameters
 - Ranks hyperparameter importance
- Other:
 - No need to build own management tool
 - Easy to share with others
 - Easy to implement on existing projects





QUESTIONS?

Background



Selected Toolset

- Python
 - Weights & Biases wandb
 - Pytorch
 - OpenCV
 - Various python modules for providing interfaces



PyTorch


Weights & Biases



Hyperparameters

Data Split

- Data is split into training, validation, and test data
 - **Training** - used to train model
 - **Validation** - used to check model accuracy during training
 - **Test** - used to test model accuracy
- Augmentations are applied after split
 - Only applied to training data
 - Not applied to validation & testing
- Split Ratios for my project
 - **Training:** 60% - 80%
 - **Validation:** 10% - 20%
 - **Test:** 10% - 20%

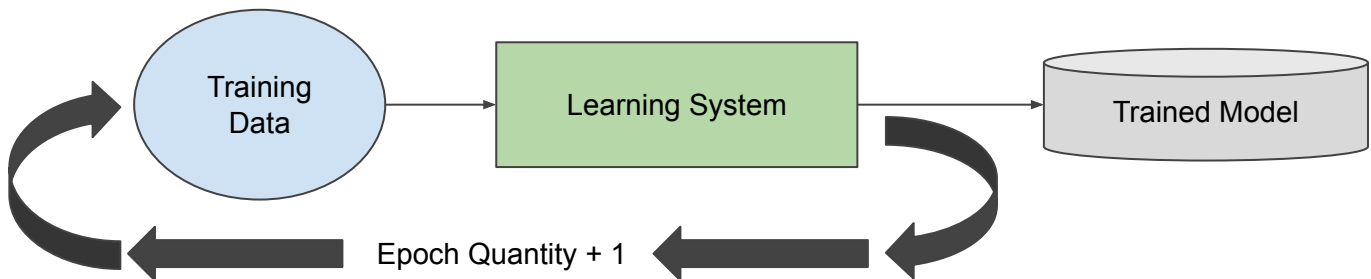




Hyperparameters

Epoch Quantity

- The number of epochs defined governs how many times the model is trained using the training data.
 - If epoch quantity is 2, then the model is trained with the training data two times.
 - Training model with the training data only one time is typically not enough to achieve a high test accuracy



Hyperparameters

Learning Rate

- Learning Rate controls how much to change the model in response to the prediction error each time the model trained.
 - Too small of learning rate results in long training process that may get stuck learning.
 - Too large of learning rate results in unstable training process where optimal learned features (weights & biases) may be skipped.
- Example
 - You're blind & walking down mountain to find bottom of valley.



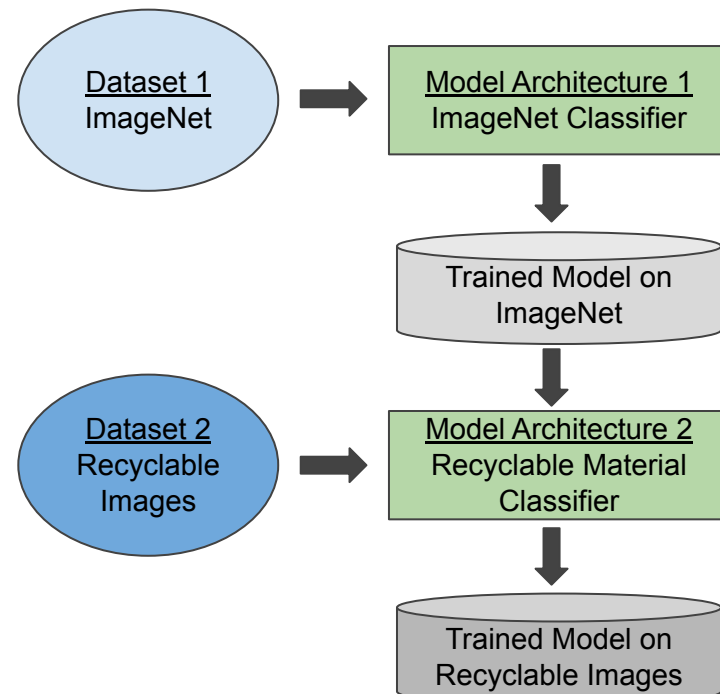
Photo: iStock



Hyperparameters

Model Architectures

- Leverage transfer learning
 - Pretrained models
 - ResNet 18, ResNet 34, ResNet 50
 - Freeze pre-trained layers to prevent back propagation
 - Define Fully Connected (FC) layer at end to learn the 6 classes
 - Customized Model
 - ResNet50-ModNet
 - Added ReLu, Dropout, FC, and LogSoftmax to last layer







Hyperparameters

Dropout

- Rate at which neurons in a layer are ignored
 - Great for preventing overfitting
 - Forces network to learn more robust features
 - Reduces training time



(a) Standard Neural Net



(b) After applying dropout.



ImageNet Dataset


ImageNet³ 2011 Fall Release (32326)

- plant, flora, plant life (4486)
- geological formation, formation (175)
- natural object (1112)
- sport, athletics (176)
- artifact, artefact (10504)
- fungus (308)
- person, individual, someone, somebody, mortal, soul (6978)
- animal, animate being, beast, brute, creature, fauna (3998)
- Misc (20400)

³ ImageNet Dataset. <http://www.image-net.org/>

Previous Work

- Top-1 vs. Top-5 Accuracy Scoring
 - **Top-1**: ground truth matching the greatest prediction
 - **Top-5**: ground truth matching one of the top 5 predictions
- ResNet⁴
 - Implements Skip Connections
 - adds outputs from previous layers to outputs of stacked layers
 - Avoids vanishing / exploding gradient
- ResNeXt⁵
 - Same as ResNet but increases number of channels from one layer to the next)
- EfficientNet⁶
 - Top-1 Accuracy: 88.61%
- Effectiveness of Data Augmentation⁷



Deep Neural Network Comparison⁸
(The size of the blobs is proportional to the number of network parameters)

⁴ Deep Residual Learning for Image Recognition. <https://arxiv.org/abs/1512.03385>

⁵ Aggregated Residual Transformations for Deep Neural Networks. <https://arxiv.org/abs/1611.05431>

⁶ EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. <https://arxiv.org/pdf/1905.11946.pdf>

⁷ The Effectiveness of Data Augmentation in Image Classification using Deep Learning. https://arxiv.org/pdf/1712.04621.pdf?source=post_page

⁸ An Analysis of Deep Neural Network Models for Practical Applications. <https://arxiv.org/abs/1605.07678>



Convolutional Neural Network (CNN)

Computer Interpretation of Image

- Image Size: 22 x 16
- Number of Pixels = 352
- Activation Number = Pixel Value



0	2	15	0	0	11	10	0	0	0	0	9	9	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	29
0	10	16	119	238	255	244	245	243	250	249	255	222	103	10	0
0	14	170	255	255	244	254	255	253	245	255	249	253	251	124	1
2	98	255	228	255	251	254	211	141	116	122	215	251	238	255	49
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36
16	229	252	254	49	12	0	0	7	7	0	70	237	252	235	62
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0
0	13	113	255	255	245	255	182	181	248	252	242	208	36	0	19
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0
0	0	0	4	58	251	255	246	254	253	255	120	11	0	1	0
0	0	4	97	255	255	255	248	252	255	244	255	182	10	0	4
0	22	206	252	246	251	241	100	24	113	255	245	255	194	9	0
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0
0	218	251	250	137	7	11	0	0	2	62	255	250	125	3	0
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4
0	18	146	250	255	247	255	255	255	249	255	240	255	129	0	5
0	0	23	113	215	255	250	248	255	255	248	248	118	14	12	0
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1
0	0	5	5	0	0	0	0	0	14	1	0	6	6	0	0


0	2	15	0	0	11	10	0	0	0	0	9	9	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	29
0	10	16	119	238	255	244	245	243	250	249	255	222	103	10	0
0	14	170	255	255	244	254	255	253	245	255	249	253	251	124	1
2	98	255	228	255	251	254	211	141	116	122	215	251	238	255	49
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36
16	229	252	254	49	12	0	0	7	7	0	70	237	252	235	62
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0
0	13	113	255	255	245	255	182	181	248	252	242	208	36	0	19
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0
0	0	0	4	58	251	255	246	254	253	255	120	11	0	1	0
0	0	4	97	255	255	255	248	252	255	244	255	182	10	0	4
0	22	206	252	246	251	241	100	24	113	255	245	255	194	9	0
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0
0	218	251	250	137	7	11	0	0	2	62	255	250	125	3	0
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4
0	18	146	250	255	247	255	255	255	249	255	240	255	129	0	5
0	0	23	113	215	255	250	248	255	255	248	248	118	14	12	0
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1
0	0	5	5	0	0	0	0	0	14	1	0	6	6	0	0



Convolutional Neural Network (CNN)

Neural Networks

- First Layer
 - **Input:** 352 pixel values
- Last Layer
 - **Output:** 10 values corresponding to number in image
- Hidden Layers
 - 16 neurons per hidden layer
 - Each input neuron has a weight
 - Each neuron recognizes feature in image

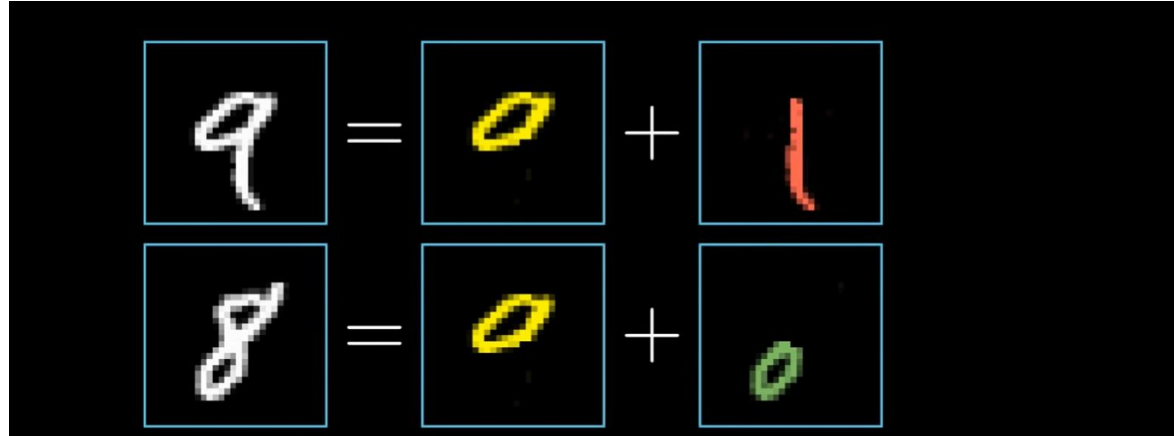





Convolutional Neural Network (CNN)

Neural Network - Hidden Layers

2nd Hidden Layer



1st Hidden Layer






Convolutional Neural Network (CNN)

Neural Network - Weights & Biases

- Connection between each neuron has a **weight** (w_n)
- **Activations** (a_n) are multiplied by weights and summed
- Activation function is applied to squish weighted sum to between 0 and 1
- Each weighted sum is assigned a **bias** (b)
 - Neuron activates if weighted sum is larger than bias
- Learnable Values = 6090
 - Weights = 6048
 $(352 \times 16) + (16 \times 16) + (16 \times 10) = 6048$
 - Biases = 42
 $16 + 16 + 10 = 42$





Convolutional Neural Network (CNN)

Neural Network - Back Propagation

1. Assign random value for weights (w_n) & biases (b)
2. Train on new image
3. Apply Cost Function to model output
 - Quantifies model predictions vs true value performance
4. Minimize Cost
 - This is called **gradient descent**
 - Step towards minimum is called **learning rate**
5. Change weight (w_n) & bias (b) values by negative gradient of cost function for each layer
6. Repeat steps 2 - 5

For more information, reference [these](#) videos.



Convolutional Neural Network (CNN)

Convolutions

1. Start with Kernel:

0	1	2
2	2	0
0	1	2

2. Slide along image pixel values
 - a. Perform an element-wise multiplication
 - b. Sum results
3. Place results into image of new pixel values

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1


12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0




Convolutional Neural Network (CNN)

Convolutions

- Convolution Parameters
 - Kernel Values
 - Kernel Size
 - Padding
 - Stride
 - Max / Min Pooling
- [More Information](#)



Stride



Padding